



# Hangman with LLMs

*Training Google's CANINE transformer to play Hangman*

*March 2024*

Vaibhav Sahu

University of Pennsylvania

✉ vaibhavs@seas.upenn.edu    🏠 vsa1920.github.io

🔗 <https://github.com/vsa1920/Hangman-with-Transformers>

# HANGMAN: A Word Guessing Game

Hangman is a popular word-guessing game:

1. Player 1 typically chooses a word that Player 2 has to guess. All the letters are concealed at the start of the game.
2. Player 2 then has to guess the letters of the hidden word, one letter at a time. All its positions are revealed if the guessed letter is in the word.
3. If a guessed letter is not in the word, Player 2 exhausts one failed try. If the player guesses the word before guessing the wrong  $N$  times, they win; otherwise, they lose.  
(typically  $N = 6$  or  $7$ )

# How hard is Hangman?

Number of letters	Optimal calling order
1	A I
2	A O E I U M B H
3	A E O I U Y H B C K
4	A E O I U Y S B F
5	S E A O I U Y H
6	E A I O U S Y
7	E I A O U S
8	E I A O U
9	E I A O U
10	E I O A U
11	E I O A D
12	E I O A F
13	I E O A
14	I E O
15	I E A
16	I E H
17	I E R
18	I E A
19	I E A
20	I E

Figure: Calling Order for the first letter according to word length  
Image courtesy - Datagenetics

# Formulating the problem

## *Hangman: Actions and States*

Applying the rules of Hangman ( $H$ ), it can be modeled as taking an action ( $a_t$ : guessing 1 of the 26 letters in the English Alphabet), given a Hangman game state ( $s_t$ : Guessed letters and the state of the hidden word)

$$H(a_t, s_t) = s_{t+1}$$

# Encoding the Hangman Game States

We want to model the game states to feed it to our Transformer model:

$$s_t = [\text{CLS}] \text{<GUESSED\_LETTERS>} [\text{SEP}] \text{<HANGMAN\_STATE>} [\text{SEP}]$$

## **Example:**

Actual word: b i r t h d a y; Guessed letters: e, a, i, s; State: \_ i \_ \_ \_ a \_

$$s_t = [\text{CLS}] \text{eias} [\text{SEP}] [\text{MASK}] \text{i} [\text{MASK}] [\text{MASK}] [\text{MASK}] [\text{MASK}] \text{a} [\text{MASK}] [\text{SEP}]$$

# Machine Learning Task

How do we model the decision that our agent takes at each step? It turns out there is more than 1 way to do so.

## 1. Multi-label Classification Task:

- Label consists of every hidden letter that hasn't been guessed.
- Model can simultaneously guess letters!
- Not too restrictive. Well defined? Hard to train?

## 2. Multi-class Classification Task:

- Guess the most likely or the most frequent letter.
- Make a distribution of the likelihood by normalizing the counts.
- A 1-1 mapping between labels and states. Easier to train?

$$[l(s_t)]_i = \frac{\sum_j \delta_j \cdot I_{x_j=x_i}}{\sum_i \sum_j \delta_j \cdot I_{x_j=x_i}}$$

where  $x_i$  is the  $i$ -th letter of the alphabet, and  $x_j$  is the letter at the  $j$ -th position of the word

# Data Generation

We have chosen the right way to label the data. The number of Hangman states in a Hangman game is huge!

$$||\Omega|| \sim 2^k \cdot \binom{26-k}{N} \cdot N!$$

$k$  : Number of unique letters in the word)

$N$ : Number of allowed wrong tries

How do we sample efficiently?

# Biased Random Sampling

## *Game Simulations*

Learn as much as possible from each simulated game!

## *Biased Random Sampling ( $p$ )*

Proposed approach:

- With a probability  $p$ , sample from one of the missing letters in the word
- The rest of the time, sample randomly from any of the letters that haven't been guessed
- $p$  is a hyperparameter to be tuned. We find  $p = 0.4$  works well!



# Model Selection: Google's CANINE

- Our states have character-level information stored in them
- Intuitively, we need a character-level language model that can model the conditional distribution of letters in English words!
- CANINE: [Clark et al., 2022]

# Training our Transformer from scratch!

The model does not benefit from transfer learning of sentence-level language modeling. This is why it's actually simpler to work with it from random initialization:

## *Training Algorithm (Pre-training)*

- Preprocess a corpus of words: Random-split to train, validate, and test
- At each epoch, for each word, the training corpus is used to perform gameplay simulation through biased sampling
- The model then learns to match the output distribution as defined before for each state, using Cross-entropy loss

## *Pre-training Results*

Our model has 50+% game-winning accuracy with 6 tries!

# Self-play fine-tuning

Our model does well. But can we improve its accuracy further? Biased sampling provides a lot of simulation states, but as the model becomes better, it needs more precise data. We need more cases where we know it does poorly.

## *Self-play fine-tuning Algorithm - motivated by [Mary Wahl, 2017]*

- Let our model navigate through the state space on its own based on its guesses
- If it guesses correctly, we proceed; otherwise, we calculate the loss
- Iteratively update the model to correct these mistakes by optimizing our loss

## *Self-play fine-tuning Results*

Our model has 63% game-winning accuracy with 6 tries and 86% with 10 tries!

# Dataset Analysis

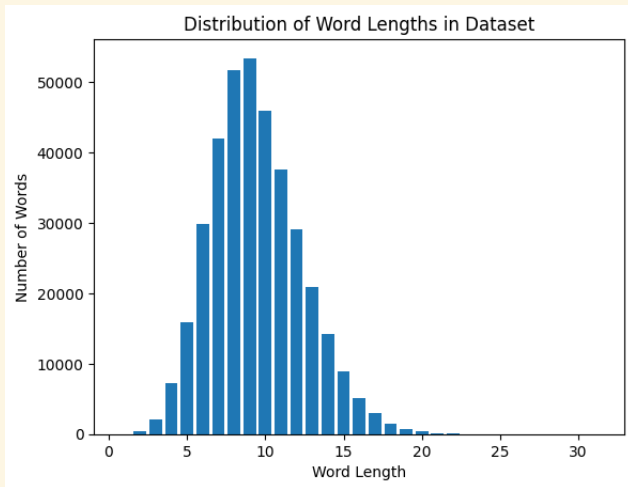


Figure: Word length distribution

# Calling Order

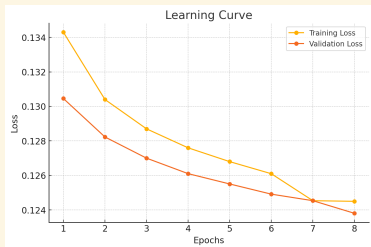
```
Length 4: ['a', 'e', 'o', 'i', 'u', 'y', 's', 'r', 'd', 't']
Length 5: ['a', 'e', 'o', 'i', 'u', 'y']
Length 6: ['e', 'a', 'o', 'i', 'u', 'y']
Length 7: ['e', 'a', 'i', 'o', 'u']
Length 8: ['e', 'a', 'i', 'o', 'u']
Length 9: ['e', 'i', 'a', 'o']
Length 10: ['e', 'i', 'o', 'a']
Length 11: ['e', 'i', 'o', 'a']
Length 12: ['e', 'i', 'o', 'a']
Length 13: ['e', 'i', 'o', 'a']
Length 14: ['i', 'e', 'o']
Length 15: ['i', 'e', 'o']
```

Figure: Calling Order of our Dataset

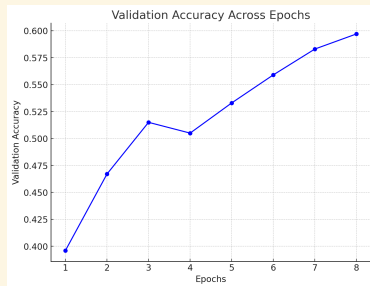
# Pre-training results

## Validation Accuracy

59 %



(a) Learning Curve

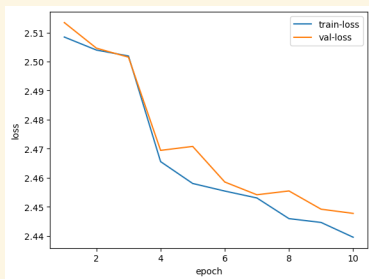


(b) Validation accuracy plot


# Self-play Finetuning

## Validation Accuracy

63.5 %



(a) Learning Curve



(b) Accuracy plot for self-play

# Model Analysis Results

## **Test Accuracy**

63 %

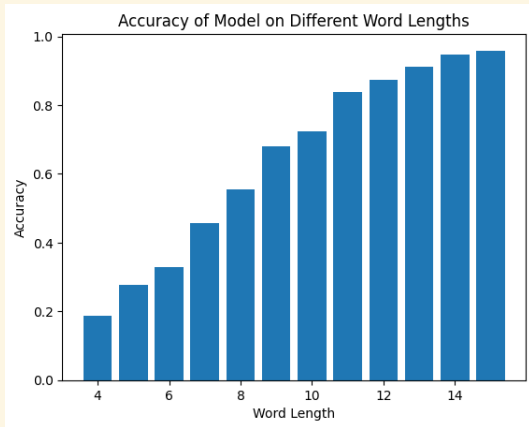


Figure: Accuracy vs word length



## Comparison with a baseline

	<b>Baseline N-gram Model</b>	<b>Modified CANINE Model</b>
6-guesses	23%	63%
10-guesses	56%	86%

**Table:** Test accuracy comparison between Baseline N-gram Model and Modified CANINE Model

# Future Scope and Improvements

## *Future Scope*

We have successfully trained our Hangman agent. What other cool downstream tasks can our trained model be used for?

1. Our model can be fine-tuned to be a great spell-checker and auto-complete! It can identify if a spelling is correct or not and even provide suggestions to auto-complete.
2. Can be an interesting decrypting tool for simple decrypting tasks!
3. Can help in speech recognition in filling missing partial phonemes.

# Future Scope and Improvements

## *Improvements*

How can we improve our model performance, reduce latency, and speed it up?

- Our model is very slow at backprops when working with single samples at a time. If we could batch them together, the fine-tuning process and iteration would achieve significant speedup.
- Parameter Efficient Fine-tuning (PEFT) is another idea worth exploring to speed up fine-tuning our model. Backpropagation would be faster even when working one sample at a time.
- More hyperparameter tuning: We have been very conservative in using Hyperparameter values, and some good enough values have worked. However, we have parameters that still require more analysis. We have a Bias rate of  $p$  and the training and fine-tuning hyperparameters such as  $\eta$  and  $\lambda$ .

# References

- Clark, J. H., Garrette, D., Turc, I., and Wieting, J. (2022).  
Canine: Pre-training an efficient tokenization-free encoder for  
language representation.  
*Transactions of the Association for Computational Linguistics*,  
10:73–91.
- Mary Wahl, Shaheen Gauher, F. B. U. K. Z. (2017).  
Microsoft azure: Play hangman with cntk.